

# CS441\_HW4\_Solution

November 3, 2025

## 0.1 CS441: Applied ML - HW 4

### 0.1.1 Part 1: Model Complexity and Tree-based Regressors

```
[ ]: import numpy as np
from google.colab import drive
%matplotlib inline
from matplotlib import pyplot as plt

# load data (modify to match your data directory or comment)
def load_temp_data():
    drive.mount('/content/drive')
    datadir = "/content/drive/MyDrive/CS441/hw1/"
    T = np.load(datadir + 'temperature_data.npz')
    x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val, \
    dates_test, feature_to_city, feature_to_day = \
    T['x_train'], T['y_train'], T['x_val'], T['y_val'], T['x_test'], T['y_test'], \
    T['dates_train'], T['dates_val'], T['dates_test'], T['feature_to_city'], \
    T['feature_to_day']
    return (x_train, y_train, x_val, y_val, x_test, y_test, dates_train, \
    dates_val, dates_test, feature_to_city, feature_to_day)

# plot one data point for listed cities and target temperature
def plot_temps(x, y, cities, feature_to_city, feature_to_day, target_date):
    nc = len(cities)
    ndays = 5
    xplot = np.array([-5,-4,-3,-2,-1])
    yplot = np.zeros((nc,ndays))
    for f in np.arange(len(x)):
        for c in np.arange(nc):
            if cities[c]==feature_to_city[f]:
                yplot[feature_to_day[f]+ndays,c] = x[f]
    plt.plot(xplot,yplot)
    plt.legend(cities)
    plt.plot(0, y, 'b*', markersize=10)
    plt.title('Predict Temp for Cleveland on ' + target_date)
    plt.xlabel('Day')
    plt.ylabel('Avg Temp (C)')
```

```

plt.show()

# load data
(x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val,
dates_test, feature_to_city, feature_to_day) = load_temp_data()

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call  
drive.mount("/content/drive", force\_remount=True).

One measure of a tree's complexity is the maximum tree depth. Train tree, random forest, and boosted tree regressors on the temperature regression task, using all default parameters except:

- max\_depth={2,4,8,16,32}
- random\_state=0
- For random forest: max\_features=1/3

Measure train and val RMSE for each and plot them all on the same plot using the provided plot\_depth\_error function. You should have six lines (train/val for each model type), each with 5 data points (one for each max depth value). Include the plot and answer the analysis questions in the report.

```
[ ]: # to plot the errors
def plot_depth_error(max_depths, tree_train_err, tree_val_err, rf_train_err,
                     rf_val_err, bt_train_err, bt_val_err):
    plt.figure()
    plt.semilogx(max_depths, tree_train_err, 'r.--', label='tree train')
    plt.semilogx(max_depths, tree_val_err, 'r.-', label='tree val')
    plt.semilogx(max_depths, rf_train_err, 'g.--', label='RF train')
    plt.semilogx(max_depths, rf_val_err, 'g.-', label='RF val')
    plt.semilogx(max_depths, bt_train_err, 'b.--', label='BT train')
    plt.semilogx(max_depths, bt_val_err, 'b.-', label='BT val')
    plt.ylabel('RMSE Error')
    plt.xlabel('Max Tree Depth')
    plt.xticks(max_depths, max_depths)
    plt.legend()
    plt.rcParams.update({'font.size': 20})
    plt.show()
```

```
[ ]: from sklearn import tree
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_squared_error

max_depths = [2,4,8,16,32]

# usage examples
# model = DecisionTreeRegressor(random_state=0, max_depth=max_depth)
```

```

# model = RandomForestRegressor(random_state=0, max_depth=max_depth, max_features=1/3)
# model = GradientBoostingRegressor(random_state=0, max_depth=max_depth)

dtr_trainerror = []
dtr_valerror = []
rfr_trainerror = []
rfr_valerror = []
gbr_trainerror = []
gbr_valerror = []

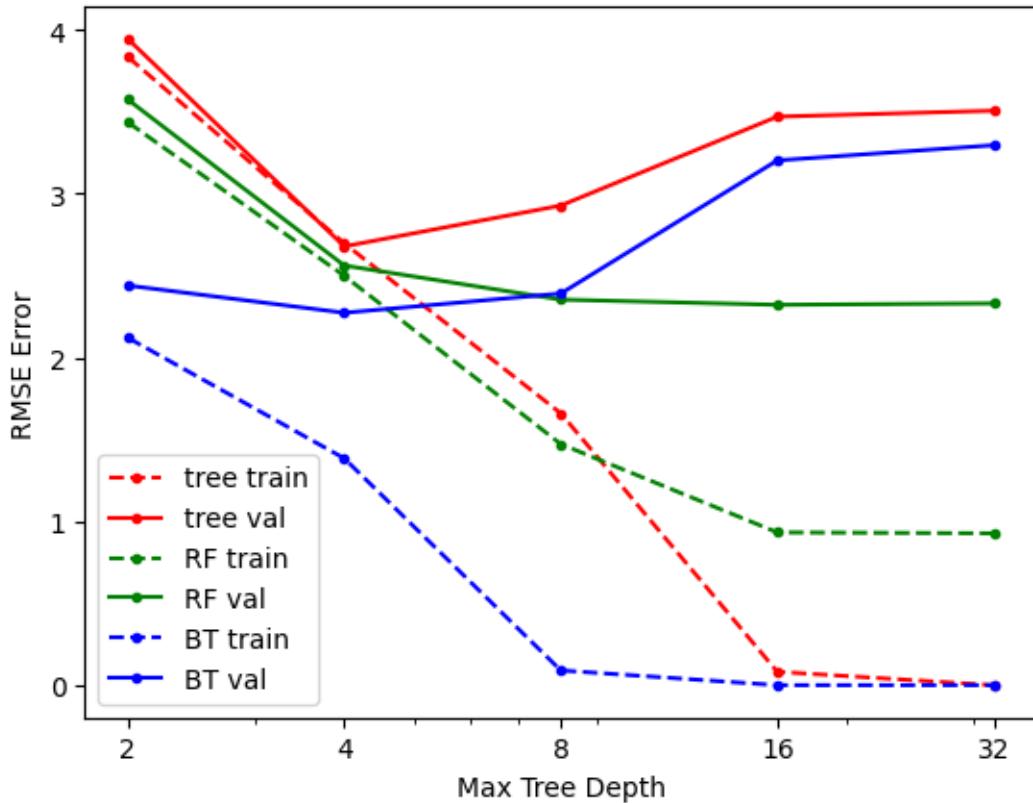
for max_depth in max_depths:
    # DecisionTreeRegressor
    model_dtr = DecisionTreeRegressor(random_state = 0, max_depth= max_depth)
    model_dtr.fit(x_train,y_train)
    dtr_train_ypredict = model_dtr.predict(x_train)
    dtr_val_ypredict = model_dtr.predict(x_val)
    dtr_trainerror.append(np.sqrt(mean_squared_error(y_train,dtr_train_ypredict)))
    dtr_valerror.append(np.sqrt(mean_squared_error(y_val,dtr_val_ypredict)))

    # RandomForestRegressor
    model_rfr = RandomForestRegressor(random_state = 0, max_depth= max_depth)
    model_rfr.fit(x_train,y_train)
    rfr_train_ypredict = model_rfr.predict(x_train)
    rfr_val_ypredict = model_rfr.predict(x_val)
    rfr_trainerror.append(np.sqrt(mean_squared_error(y_train,rfr_train_ypredict)))
    rfr_valerror.append(np.sqrt(mean_squared_error(y_val,rfr_val_ypredict)))

    # GradientBoostingRegressor
    model_gbr = GradientBoostingRegressor(random_state = 0, max_depth= max_depth)
    model_gbr.fit(x_train,y_train)
    gbr_train_ypredict = model_gbr.predict(x_train)
    gbr_val_ypredict = model_gbr.predict(x_val)
    gbr_trainerror.append(np.sqrt(mean_squared_error(y_train,gbr_train_ypredict)))
    gbr_valerror.append(np.sqrt(mean_squared_error(y_val,gbr_val_ypredict)))

plot_depth_error(max_depths, dtr_trainerror, dtr_valerror, rfr_trainerror, rfr_valerror, gbr_trainerror, gbr_valerror)

```



### 0.1.2 Part 2: MLPs with MNIST

For this part, you will want to use a GPU to improve runtime. Google Colab provides limited free GPU acceleration to all users. Go to Runtime and change Runtime Type to GPU. This will reset your compute node, so do it before starting to run other cells.

See [Tips](#) for detailed guidance on this problem.

First, use PyTorch to implement a Multilayer Perceptron network with one hidden layer (size 64) with ReLU activation. Set the network to minimize cross-entropy loss, which is the negative log probability of the training labels given the training features. This objective function takes unnormalized logits as inputs.

*Do not use MLP in sklearn for this HW - use Torch.*

```
[ ]: # initialization code
import numpy as np
from keras.datasets import mnist
%matplotlib inline
from matplotlib import pyplot as plt
from scipy import stats
import torch
import torch.nn as nn
```

```

def load_mnist():
    """
    Loads, reshapes, and normalizes the data
    """

    (x_train, y_train), (x_test, y_test) = mnist.load_data() # loads MNIST data
    x_train = np.reshape(x_train, (len(x_train), 28*28)) # reformat to 768-d
    ↪vectors
    x_test = np.reshape(x_test, (len(x_test), 28*28))
    maxval = x_train.max()
    x_train = x_train/maxval # normalize values to range from 0 to 1
    x_test = x_test/maxval
    return (x_train, y_train), (x_test, y_test)

def display_mnist(x, subplot_rows=1, subplot_cols=1):
    """
    Displays one or more examples in a row or a grid
    """

    if subplot_rows>1 or subplot_cols>1:
        fig, ax = plt.subplots(subplot_rows, subplot_cols, figsize=(15,15))
        for i in np.arange(len(x)):
            ax[i].imshow(np.reshape(x[i], (28,28)), cmap='gray')
            ax[i].axis('off')
    else:
        plt.imshow(np.reshape(x, (28,28)), cmap='gray')
        plt.axis('off')
    plt.show()

```

```
[ ]: # Sets device to "cuda" if a GPU is available (in Colabs, enable GPU by
    ↪Edit->Notebook Settings-->Hardware Accelerator=GPU)
device = "cuda" if torch.cuda.is_available() else 'cpu'
print(device) # make sure you're using GPU instance
```

cuda

**2a** Using the train/val split provided in the starter code, train your network for 100 epochs with learning rates of 0.01, 0.1, and 1. Use a batch size of 256 and the SGD optimizer. After each epoch, record the mean training and validation loss and compute the validation error of the final model. The mean validation loss should be computed after the epoch is complete. The mean training loss can either be computed after the epoch is complete, or, for efficiency, computed using the losses accumulated during the training of the epoch. Plot the training and validation losses using the display\_error\_curves function.

```
[ ]: (x_train, y_train), (x_test, y_test) = load_mnist()

# create train/val split
ntrain = 50000
```

```
x_val = x_train[ntrain:].copy()
y_val = y_train[ntrain:].copy()
x_train = x_train[:ntrain]
y_train = y_train[:ntrain]
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/mnist.npz
11490434/11490434           2s
0us/step
```

```
[ ]: def display_error_curves(training_losses, validation_losses):
    """
    Plots the training and validation loss curves
    training_losses and validation_losses should be lists or arrays of the same_
    length
    """
    num_epochs = len(training_losses)

    plt.plot(range(num_epochs), training_losses, label="Training Loss")
    plt.plot(range(num_epochs), validation_losses, label="Validation Loss")

    # Add in a title and axes labels
    plt.title('Training and Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')

    # Display the plot
    plt.legend(loc='best')
    plt.show()
```

```
[ ]: # Define the model
class MLP(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(MLP, self).__init__()
        # Needs code here
        self.model = nn.Sequential(
            nn.Linear(input_size,hidden_size),
            nn.ReLU(),
            nn.Linear(hidden_size, output_size)
        )

    def forward(self, x):
        # Needs code here
        return self.model(x)
```

```
[ ]: print(x_train.shape)
```

(50000, 784)

```
[ ]: # This is a possible function definition for training MLP, but feel free to
    ↪change it
# You may also want to create helper functions, e.g. for computing loss or
    ↪prediction
def train_MLP_mnist(train_loader, val_loader, lr=1e-1, num_epochs=100):
    """
    Train a MLP
    Input: train_loader and val_loader are dataloaders for the training and
    val data, respectively. lr is the learning rate, and the network will
    be trained for num_epochs epochs.
    Output: return a trained MLP
    """
    # TODO: fill in all code

    input_size = x_train.shape[1]
    hidden_size = 64
    output_size = 10

    # Instantiate the model
    mlp = MLP(input_size,hidden_size,output_size).to(device)
    criterion = torch.nn.CrossEntropyLoss()
    optimizer = torch.optim.SGD(mlp.parameters(),lr=lr) # 0.01, 0.1, 1
    train_loss_list = []
    val_loss_list = []
    val_loss_epochmin = 0

    # Train the model, compute and store train/val loss at each epoch

    epochs = 100
    for epoch in range(epochs):
        mlp.train()
        running_loss = 0
        n = 0
        for inputs,targets in train_loader:
            inputs = inputs.to(device)
            targets = targets.to(device)

            outputs = mlp(inputs)
            loss = criterion(outputs,targets)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            # loss.item() batch
```

```

        running_loss += loss.item()*len(targets) # last batch 256
        n += len(targets) # sample numbers of all epoch

    train_loss_list.append(running_loss/n) # loss of an epoch

    mlp.eval()
    val_loss, val_error = evaluate_MLP(mlp, val_loader)
    val_loss_list.append(val_loss)

    val_loss_epochmin = np.argmin(val_loss_list)
    print(f"Epoch [{epoch+1}/{num_epochs}], Train Loss: {train_loss_list[-1]:.4f}, Val Loss: {val_loss_list[-1]:.4f}")
    print("Best_Epoch:", val_loss_epochmin)

    # Display Loss Curves
    display_error_curves(train_loss_list, val_loss_list)

    return mlp

def evaluate_MLP(mlp, loader):
    ''' Computes loss and error rate given your mlp model and data loader'''
    N = 0
    acc = 0
    loss = 0
    loss_function = torch.nn.CrossEntropyLoss()
    with torch.set_grad_enabled(False):
        for i, data in enumerate(loader, 0):

            # Get inputs
            inputs, targets = data
            N += len(targets)

            # Perform forward pass
            outputs = mlp(inputs.to(device))

            # Compute sum of correct labels
            y_pred = np.argmax(outputs.cpu().numpy(), axis=1)
            y_gt = targets.numpy()
            acc += np.sum(y_pred==y_gt)

            # Compute loss
            loss += loss_function(outputs, targets.to(device)).item()*len(targets)

    loss /= N
    acc /= N

```

```

    return loss, 1-acc

[ ]: from torch.utils.data import DataLoader, TensorDataset
# Code for running experiments

print(device) # make sure you're using GPU instance
torch.manual_seed(0) # to avoid randomness, but if you wanted to create an
#ensemble, you should not use a manual seed

# TODO (set up dataloaders, and call training function)
x_train_tensor = torch.tensor(x_train, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train, dtype=torch.long)

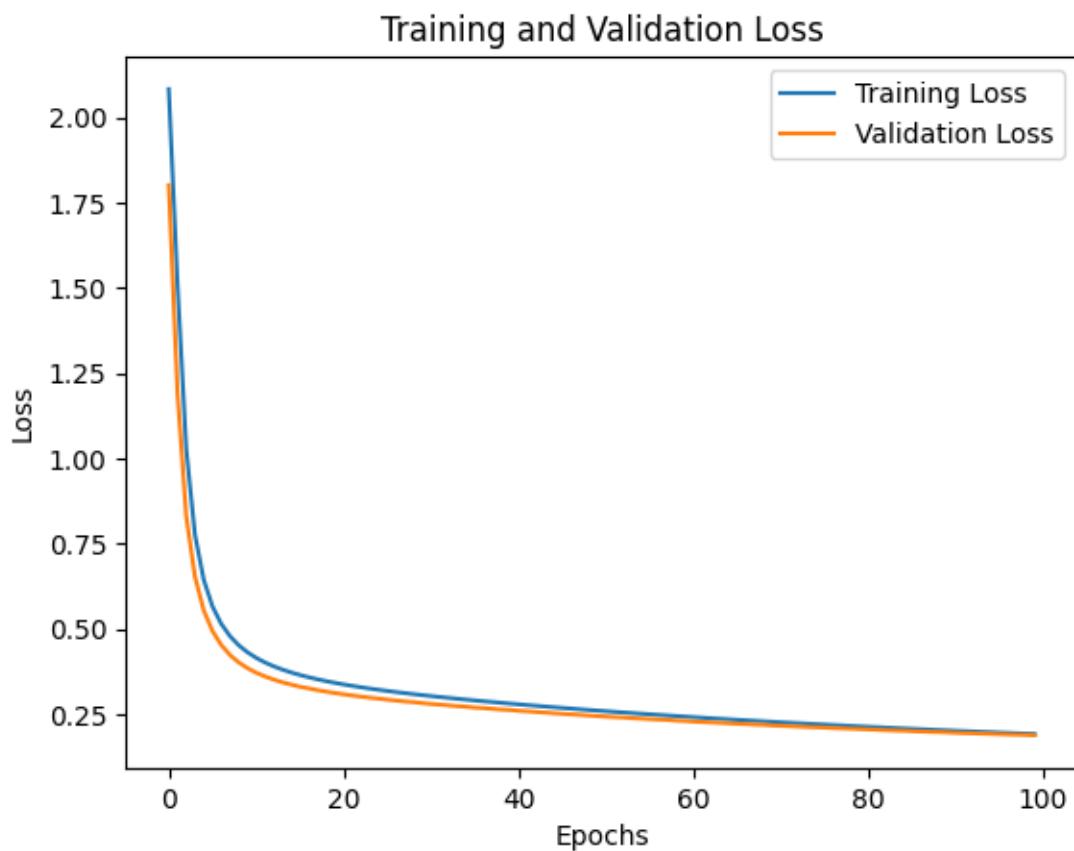
x_val_tensor = torch.tensor(x_val, dtype=torch.float32)
y_val_tensor = torch.tensor(y_val, dtype=torch.long)

batch_size = 256
train_loader = 
    DataLoader(TensorDataset(x_train_tensor,y_train_tensor),batch_size = 256,shuffle=True)
val_loader = DataLoader(TensorDataset(x_val_tensor,y_val_tensor),batch_size = 256,shuffle=False)
lrs = [0.01,0.1,1]

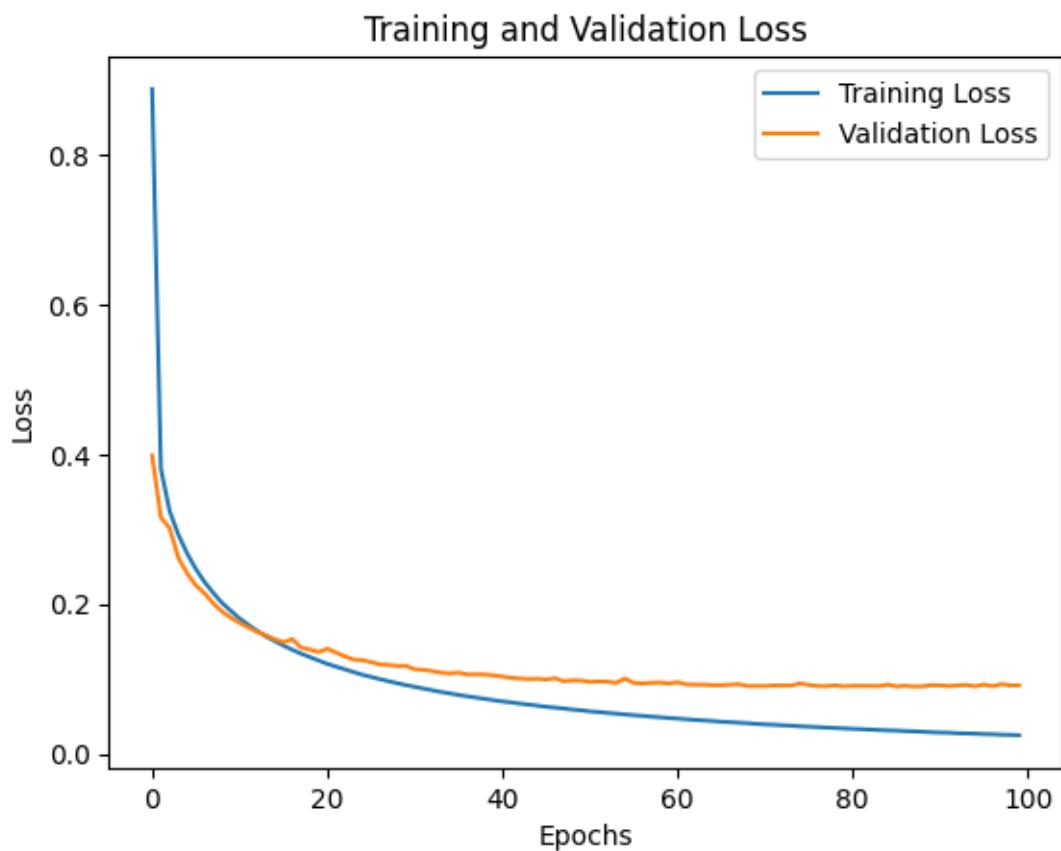
for lr in lrs:
    model = train_MLP_mnist(train_loader, val_loader, lr=lr)
    evaluate_MLP(model,val_loader)

```

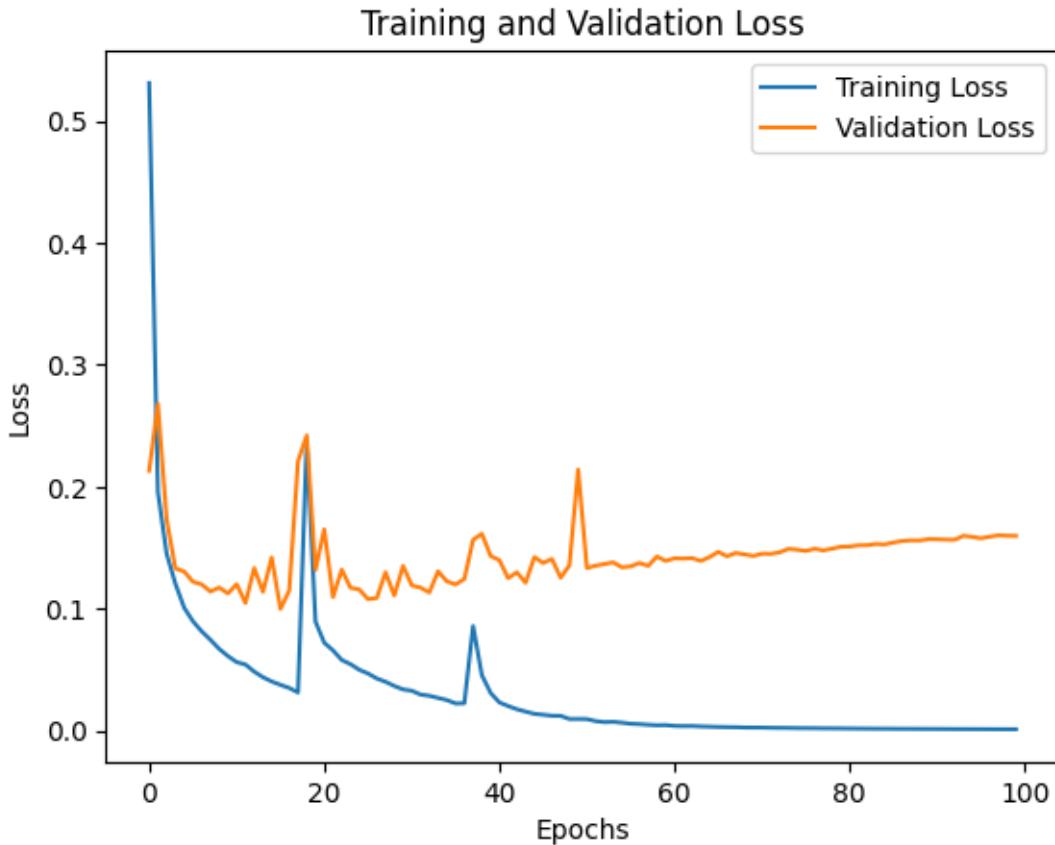
cuda  
Epoch [100/100], Train Loss: 0.1909, Val Loss: 0.1873  
Best\_Epoch: 99



Epoch [100/100], Train Loss: 0.0250, Val Loss: 0.0916  
Best\_Epoch: 87



Epoch [100/100], Train Loss: 0.0011, Val Loss: 0.1597  
Best\_Epoch: 15



**2b** Based on the loss curves, select the learning rate and number of epochs that minimizes the validation loss. Retrain that model (if it's not stored), and report training loss, validation loss, training error, validation error, and test error.

```
[ ]: # TO DO (retrain if needed, and evaluate model on train, val, and test sets)
best_lr = 0.1
best_epoch = 87

best_model = MLP(x_train.shape[1], 64, 10).to(device)
criterion = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(best_model.parameters(), lr=best_lr)

for epoch in range(best_epoch):
    best_model.train()
    for inputs,targets in train_loader:
        inputs = inputs.to(device)
        targets = targets.to(device)

        optimizer.zero_grad()
```

```

outputs = best_model(inputs)
loss = criterion(outputs, targets)
loss.backward()

optimizer.step()

x_test_tensor = torch.tensor(x_test, dtype=torch.float32)
y_test_tensor = torch.tensor(y_test, dtype=torch.long)
test_loader = DataLoader(TensorDataset(x_test_tensor,y_test_tensor),batch_size=256,shuffle=False)

train_loss, train_error = evaluate_MLP(best_model, train_loader)
val_loss, val_error = evaluate_MLP(best_model, val_loader)
_, test_error = evaluate_MLP(best_model, test_loader)

print(f"Training Loss: {train_loss:.1E}, Training Error: {train_error:.4f}")
print(f"Validation Loss: {val_loss:.1E}, Validation Error: {val_error:.4f}")
print(f"Test Error: {test_error:.4f}")

```

Training Loss: 2.9E-02, Training Error: 0.0061  
 Validation Loss: 8.8E-02, Validation Error: 0.0239  
 Test Error: 0.0257

## 0.2 Part 3: Predicting Penguin Species

Include all your code for part 3 in this section.

```
[ ]: import numpy as np
from google.colab import drive
%matplotlib inline
from matplotlib import pyplot as plt
import pandas as pd
import seaborn as sns
#styling preferences for sns
sns.set_style('whitegrid')
sns.set_context('poster')
drive.mount('/content/gdrive/')
datadir = "/content/gdrive/MyDrive/CS441/hw4/" # TO DO: modify this to your
#directory
df_penguins = pd.read_csv(datadir + 'penguins_size.csv')
df_penguins.head(10)

# convert features with multiple string values to binary features so they can
# be used by sklearn
def get_penguin_xy(df_penguins):
    data = np.array(df_penguins[['island', 'culmen_length_mm', 'culmen_depth_mm',
                                'flipper_length_mm', 'body_mass_g', 'sex']])

```

```

y = df_penguins['species']
ui = np.unique(data[:,0]) # unique island
us = np.unique(data[:, -1]) # unique sex
X = np.zeros((len(y), 10))
for i in range(len(y)):
    f = 0
    for j in range(len(ui)):
        if data[i, f]==ui[j]:
            X[i, f+j] = 1
    f = f + len(ui)
    X[i, f:(f+4)] = data[i, 1:5]
    f=f+4
    for j in range(len(us)):
        if data[i, 5]==us[j]:
            X[i, f+j] = 1
feature_names = ['island_biscoe', 'island_dream', 'island_torgersen', ↴
↳ 'culmen_length_mm', 'culmen_depth_mm', 'flipper_length_mm', 'body_mass_g', ↴
↳ 'sex_female', 'sex_male', 'sex_unknown']
X = pd.DataFrame(X, columns=feature_names)
return(X, y, feature_names, np.unique(y))

```

Drive already mounted at /content/gdrive/; to attempt to forcibly remount, call drive.mount("/content/gdrive/", force\_remount=True).

**3a** Spend some time to visualize different pairs of features and their relationships to the species. We've done one for you. Include in your report at least two other visualizations.

```
[ ]: def plot_scatter(feature1, feature2):
    """
    Provide names of two features to create a scatterplot of them
    E.g. plot_scatter('culmen_length_mm', 'culmen_depth_mm')
    Possible features: 'culmen_length_mm', 'culmen_depth_mm', ↴
    ↳ 'flipper_length_mm', 'body_mass_g'
    """

    palette = ["red", "blue", "orange"]

    sns.scatterplot(data=df_penguins, x = feature1, y = feature2,
                    hue = 'species', palette=palette, alpha=0.8)
    # Doc: https://seaborn.pydata.org/generated/seaborn.scatterplot.html

    plt.xlabel(feature1, fontsize=14)
    plt.ylabel(feature2, fontsize=14)
    plt.title(feature1 + ' vs ' + feature2, fontsize=20)
    plt.legend(bbox_to_anchor=(1.0, 1.0), loc='upper left')
    plt.show()
```

```

# TO DO call plot_scatter with different feature pairs to create some visualizations

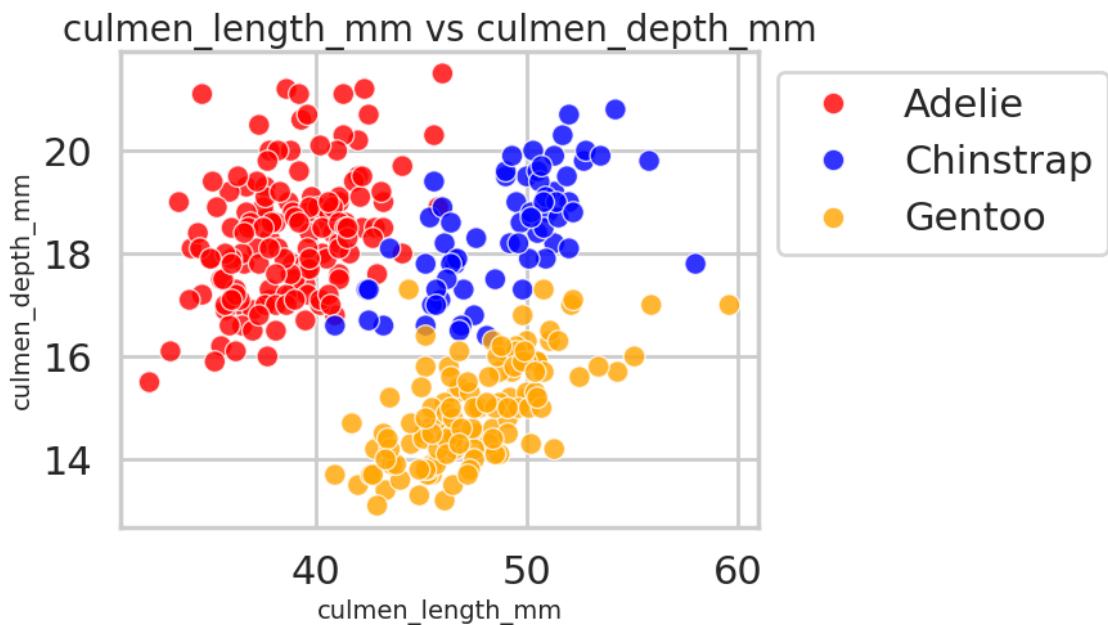
plot_scatter('culmen_length_mm', 'culmen_depth_mm')

# Visualization 1
plot_scatter('culmen_depth_mm', 'flipper_length_mm')

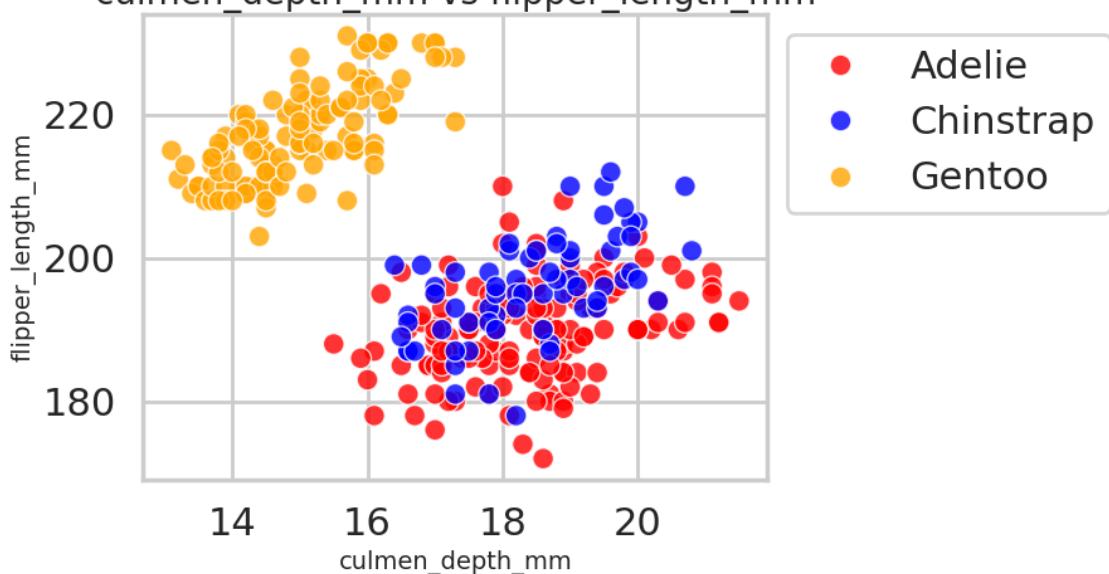
# Visualization 2
plot_scatter('flipper_length_mm', 'culmen_length_mm')

# Visualization 3
plot_scatter('flipper_length_mm', 'body_mass_g')

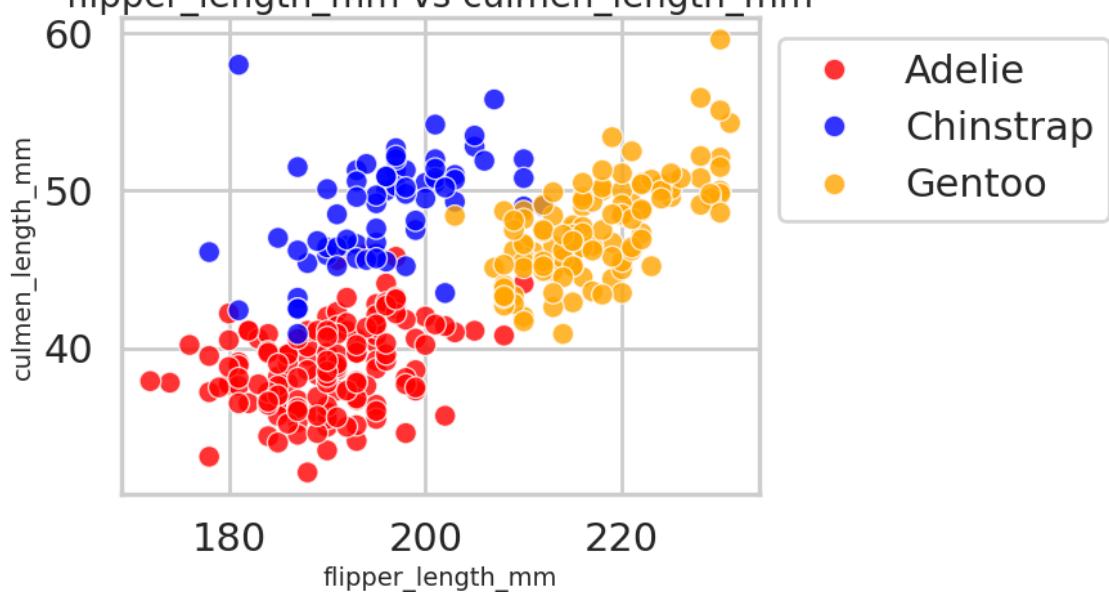
```

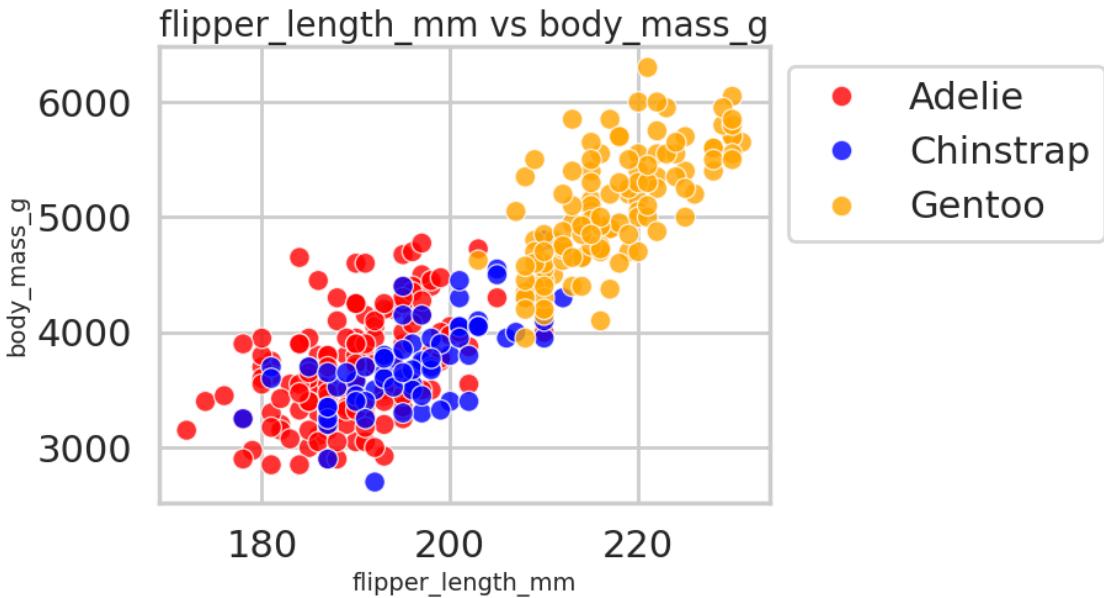


culmen\_depth\_mm vs flipper\_length\_mm



flipper\_length\_mm vs culmen\_length\_mm





**3b** Suppose you want to be able to identify the Gentoo species with a simple rule with very high accuracy. Use a decision tree classifier to figure out such a rule that has only two checks (e.g. “mass greater than 4000 g, and culmen length less than 40 mm is Gentoo; otherwise, not”). You can use the library DecisionTreeClassifier with either ‘gini’ or ‘entropy’ criterion. Use sklearn.tree.plot\_tree with feature\_names and class\_names arguments to visualize the decision tree. Include the tree that you used to find the rule in your report and the rule.

```
[ ]: # TO DO (Train a short tree to identify a good rule, plot the tree, report the rule and its precision/recall in your report)
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import classification_report, precision_score, recall_score
import matplotlib.pyplot as plt

X, y, feature_names, class_names = get_penguin_xy(df_penguins)
print(feature_names, class_names)

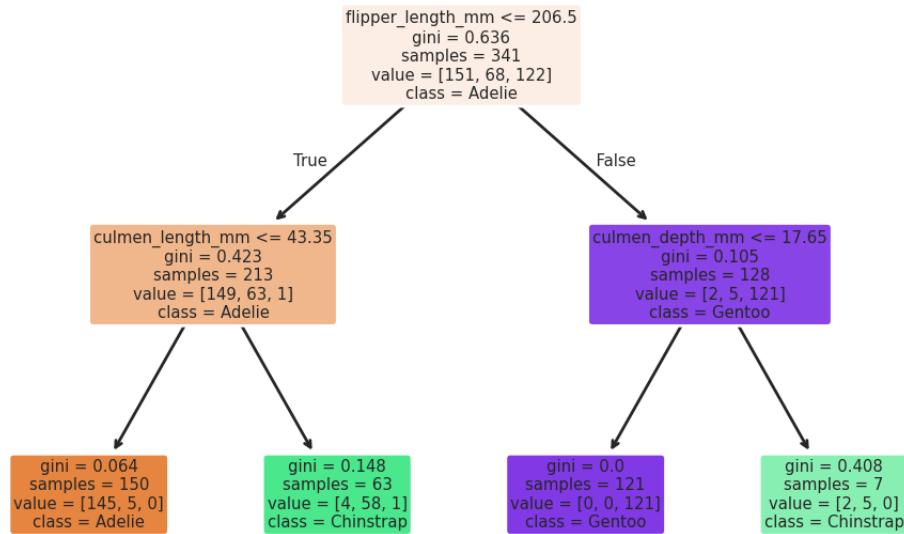
model = DecisionTreeClassifier(max_depth = 2, criterion='gini', random_state=0)
model.fit(X,y)

plt.figure(figsize=(12, 8))
plot_tree(model, feature_names=feature_names, class_names=class_names, filled=True, rounded=True)
plt.title("Decision Tree for Identifying Gentoo Penguins (max_depth=2)")
plt.show()

y_pred = model.predict(X)
print(classification_report(y, y_pred))
```

```
['island_biscoe', 'island_dream', 'island_torgersen', 'culmen_length_mm',
'culmen_depth_mm', 'flipper_length_mm', 'body_mass_g', 'sex_female', 'sex_male',
'sex_unknown'] ['Adelie' 'Chinstrap' 'Gentoo']
```

## Decision Tree for Identifying Gentoo Penguins (max\_depth=2)



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Adelie       | 0.97      | 0.96   | 0.96     | 151     |
| Chinstrap    | 0.90      | 0.93   | 0.91     | 68      |
| Gentoo       | 1.00      | 0.99   | 1.00     | 122     |
| accuracy     |           |        | 0.96     | 341     |
| macro avg    | 0.96      | 0.96   | 0.96     | 341     |
| weighted avg | 0.97      | 0.96   | 0.97     | 341     |

**3c** Use any method at your disposal to achieve maximum 5-fold cross-validation accuracy on this problem. To keep it simple, we will use `sklearn.model_selection` to perform the cross-validation for us. Report your model design and 5-fold accuracy. It is possible to get more than 99% accuracy.

```
[ ]: # design a classification model, import libraries as needed
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier

X, y, feature_names, class_names = get_penguin_xy(df_penguins)
```

```
# TO DO -- choose some model and fit the data
model = RandomForestClassifier( n_estimators=200,max_depth=None,random_state=0 )

scores = cross_val_score(model, np.array(X), np.array(y), cv=5)
print('CV Accuracy: {}'.format(scores.mean()))
```

CV Accuracy: 0.9911764705882353

### 0.3 Part 4: Stretch Goals

Include any new code needed for Part 4 here

[ ]: # TO DO (optional)

```
# from https://gist.github.com/jonathanagustin/b67b97ef12c53a8dec27b343dca4abba
# install can take a minute

import os
# @title Convert Notebook to PDF. Save Notebook to given directory
NOTEBOOKS_DIR = "/content/drive/MyDrive/CS441/hw4" # @param {type:"string"}
NOTEBOOK_NAME = "CS441_HW4_Solution.ipynb" # @param {type:"string"}
#-----#
from google.colab import drive
drive.mount("/content/drive/", force_remount=True)
NOTEBOOK_PATH = f"{NOTEBOOKS_DIR}/{NOTEBOOK_NAME}"
assert os.path.exists(NOTEBOOK_PATH), f"NOTEBOOK NOT FOUND: {NOTEBOOK_PATH}"
!apt install -y texlive-xetex texlive-fonts-recommended texlive-plain-generic > /dev/null 2>&1
!jupyter nbconvert "$NOTEBOOK_PATH" --to pdf > /dev/null 2>&1
NOTEBOOK_PDF = NOTEBOOK_PATH.rsplit('.', 1)[0] + '.pdf'
assert os.path.exists(NOTEBOOK_PDF), f"ERROR MAKING PDF: {NOTEBOOK_PDF}"
print(f"PDF CREATED: {NOTEBOOK_PDF}")
```

Mounted at /content/drive/